

Engineering Rank and Select Queries on Wavelet Trees

Roland Larsen Pedersen

Datalogi, Aarhus Universitet

Thesis defence

June 25, 2015

- 1 What is a Wavelet Tree?
 - Definitions
 - Constructing the Wavelet Tree
- 2 Queries
 - Rank
 - Select
- 3 Applications
 - Information Retrieval and Compression
- 4 Experiments and Results
- 5 Conclusion


What is a wavelet tree?

Wavelet Tree: Definitions

- Balanced binary tree.
- Stores a *sequence* $S[1, n] = c_1 c_2 c_3 \dots c_n$ of *symbols* $c_i \in \Sigma$, where $\Sigma = [1 \dots \sigma]$ is the *alphabet* of S .
- Height $h = \lceil \log \sigma \rceil$.
- $2\sigma - 1$ nodes
- Construction time: $O(n \log \sigma)$
- Memory usage: $O(n \log \sigma + \sigma \cdot ws)$ bits.
 - ws er wordsize og $ws \geq \log n$.

Constructing the Wavelet Tree

- The wavelet tree is constructed recursively.
- Each node calculates the middle character of Σ and uses it to set the bits in the bitmap and split S in two substrings S_{left} and S_{right} .



$S = \text{adsfadaadsfaads}, \Sigma = \text{adfs}$


Queries

Wavelet Tree: Queries

- The wavelet tree supports three queries:
 - **Access(p)**: Return the character c at position p in sequence S .
 - Running time: $O(n \log \sigma)$.
 - Can be reduced to $O(\log \sigma)$ using a minimal amount of extra space.
 - **Rank(c, p)**: Return the number of occurrences of character c in S up to position p .
 - Running time: $O(n \log \sigma)$.
 - Can be reduced to $O(\log \sigma)$ using a minimal amount of extra space.
 - **Select(c, o)**: Return the position of the o th occurrence of character c in S .
 - Running time: $O(n \log \sigma)$
 - Can be reduced to $O(\log \sigma)$ using a minimal amount of extra space.

Example: Access

Query = Access($p=7$).




$S = \text{adsfadaadsfaads}$, $\Sigma = \text{adfs}$

Rank on a Wavelet Tree

$$\square = p$$

Query: Rank(c='a', p=10), Result: o = 3

0 = bit representing 'a'




Select on a Wavelet Tree

 = p

Query: $\text{Select}(c='a', o=3)$, Result: $p = 9$

0 = bit representing 'a'



Applications

Information Retrieval and Compression

- Information Retrieval

- Positional inverted index.

- For each word: Return positions of occurrences.
 - $S = \text{words, can, consist, of, several, words.}$
 - Position of $word = \text{Select}(\text{words}, 2)$
 - Word at position $p = \text{Access}(p)$.

- Document retrieval.

- Return what document a word appears in.
 - Example: $S = \$\text{dasdfsd}\$ \text{fadfadfa}\$ \text{afadfadfadf}\$ \text{dfasgsdag}$
 - Position p in $S = \text{Select}_c(f, 7)$.
 - Document number $dn = \text{Rank}_s(\$, p)$.
 - Position within document $= p - \text{Select}(\$, dn)$

- Range Quantile Query.

- Return the k th smallest number within a subsequence of a given sequence of elements.

- Compression

- Zero-order entropy compression (H_0)

- The amount of compression it is possible to achieve when looking at each symbol in a sequence individually.

- Higher-order entropy compression (H_k)


- The amount of compression that can be achieved when looking at a symbol and its context.

- $H_k \leq H_0 \leq \log \sigma$.


H_0 Compression: Run-length encoding

- Run-length Encoding (RLE)
- Example: $RLE(\text{aaaaabbbaacccccaaaaa}) = \text{a5,b3,a2,c5,a5}$.
- Binary example: $RLE(00000000001111100000) = 10, 5, 5$
- Query by reversing RLE. It takes linear time $O(n)$ to reverse. Rank and select query time becomes $O(2n \log \sigma) = O(n \log \sigma)$
- Because it looks at each symbol individually it achieves H_0 compression.

Run-length encoded Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



(a) Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



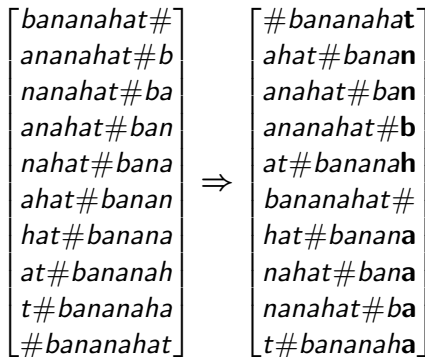
(b) RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$

Compression: Burrows-Wheeler transform

- BWT permutes the order of the characters.
- Group characters based on context.
- As a result it groups symbols more which improves the effect of Run-length encoding
- BWT is reversible
- Combined with RLE Wavelet Tree it achieves H_k compression.


BWT example

$S = \text{bananahat.}$




$BWT(S) = \text{tnnbhaaaa.}$

RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



(a) RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$




(b) BWT RLE Wavelet Tree on string *tnnbhaaaa* with alphabet $\Sigma = abhnt$


Huffman shaped wavelet tree

- Use Huffman codes of symbols to shape the tree.
- Most frequent symbols at the top of the tree.
- Least frequent symbols at the bottom of the tree.
- Best using non-uniformly distributed data like a natural language text.

Huffman Shaped Wavelet Tree: Example



(a) Balanced Wavelet tree: 39 bits



(b) Huffman-shaped wavelet tree: 22 bits

Huffman Shaped WT: Space complexity

- Balanced version: $n \log \sigma + o(n \log \sigma) + O(\sigma \cdot ws)$ bits
- Huffman-shaped: $n(H_0(S) + 1) + o(n(H_0(S) + 1)) + O(\sigma \cdot ws)$ bits.
 - From [Efficient Compressed Wavelet Trees over Large Alphabets by Navarro et al.]

Experiments and Results

Focus of experiments


- Focus on optimizing and observing the effect of hardware penalties.
 - Cache Misses.
 - Branch Mispredictions.
 - Translation Lookaside Buffer (TLB) Misses.

Experiments


- Calculate binary rank and select using popcount.
 - Popcount counts number of 1's in a binary number of size 64 bit in $O(1)$ time.
 - Reduces time spent in binary rank and select.
 - Reduces branch mispredictions and cpu cycles
- Pre-compute binary rank values in blocks.
 - Loop blocks of precomputed rank values to reduce time spent in in binary rank and select.
 - Reduces running time and cache misses
- Concatenate bitmaps and Page-align blocks.
 - Concatenate bitmaps to reduce memory and page-align blocks of precomputed ranks to reduce TLB misses.
- Pre-compute cumulative sums of rank values.
 - Remove need to linear scan blocks.
 - Enables binary rank in $O(b)$ time.

Calculate binary rank and select using popcount

- Rank: Running time $O(n \log \sigma)$.
- Select: Running time $O(n \log \sigma)$.



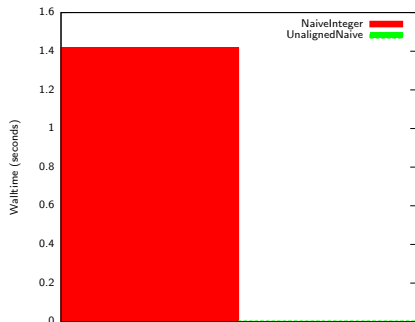
(a) Rank



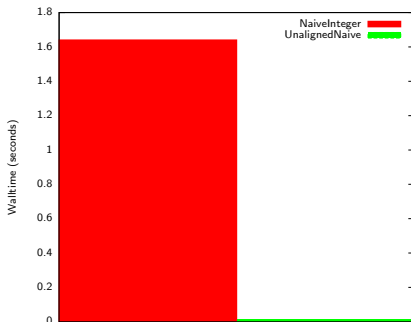
(b) Select

Pre-compute binary rank values in blocks

- Rank: Running time $O((\frac{n}{b} + b) \log \sigma)$.
- Select: Running time $O((\frac{n}{b} + b) \log \sigma)$.



(a) Rank




(b) Select


Figure : Comparison of wall time of rank and select queries between SimpleNaive not using precomputed values and UnalignedNaive using precomputed values.

Calculate binary rank and select using popcount

- Increased branch misprediction because of extra checks to handle edge cases.



(a) Rank




(b) Select


The various precomputed versions

Name	Concatenated Bitmaps	Page-aligned Blocks
Preallocated	yes	yes
UnalignedPreallocated	yes	no
Naive	no	yes
UnalignedNaive	no	no

Running time: Pre-compute binary rank values in blocks



(a) Rank: Running Time




(b) Select: Running Time


Best Block size: $\frac{1}{2}$ page size = $\frac{1}{2} * 4096$ bytes = 2048 bytes.

Memory usage: Pre-compute binary rank values in blocks

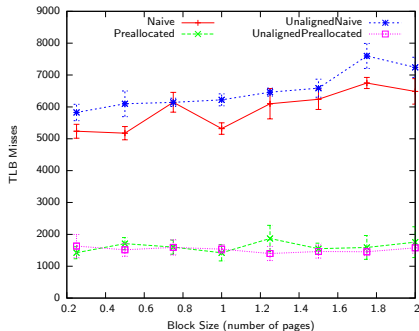
- There are $O(\frac{n}{b})$ blocks per level of the tree, and so an extra memory consumption of $O(\frac{n}{b} \log \sigma)$ words making the total memory consumption $O(n \log \sigma + (\sigma + \frac{n}{b} \log \sigma) \cdot ws)$ bits.



Page-align TLB misses



(a) Rank: TLB Misses




(b) Select: TLB Misses

- *Naive* does reduce TLB misses because of page alignment.
- Concatenated bitmaps reduces TLB misses, but page-aligning does not have much effect.


Cumulative sum

- Each block contain sum of previous blocks.
- Binary rank in $O(b)$ time in stead of $O(\frac{n}{b} + b)$ time.
- Binary search in select.
- Work per level change from $O(\frac{n}{b} + b)$ to $O(\log \frac{n}{b} + b)$. Select query total work $O((\log \frac{n}{b} + b) \log \sigma)$.
- A block size below 64 bits should not be an improvement because popcount works on words of size 64 bits.


Cumulative sum: Rank and Select running time



(a) Rank.



(b) Select.



(c) Memory Usage.

Conclusion

Conclusion: What did we learn?

- What effect hardware can have on running time and memory.
- How to do tests and how to show their results in an understandable way.
- The wavelet tree has many applications.
- The wavelet tree is great for compression of natural language texts.
- How to do literature search and how important it is.
- In general, improvements that reduced the raw amount of computations and memory accesses needed were a big improvement.
- That a simple concept can be very difficult to implement.
- Gained experience with profilers and hardware measurement tools (cachegrind, PAPI, Massif)

Conclusion: Problems and questions we faced

- Should we use uniform or non-uniform data?
- How should non-uniform data be distributed?
- How large alphabet and input size should we use?
- Debugging implementation errors in c++
- Making the implementations work
- Should we have focused on compression instead?
- PAPI produced weird memory measurements. Figuring out what was wrong took some time.
- How to avoid introducing bias in tests.

The End