

Engineering Rank and Select Queries on Wavelet Trees

Roland Larsen Pedersen

Datalogi, Aarhus Universitet

Thesis defence

June 25, 2015

Overview

1 What is a Wavelet Tree?

- Definitions
- Constructing the Wavelet Tree

2 Queries

- Rank
- Select

3 Applications

- Information Retrieval
- Compression
 - Run-length encoding
 - Burrows-Wheeler Transform
 - Huffman Shaped Wavelet tree

4 Experiments and Results

- Popcount
- Pre-compute binary rank values in blocks
- Concatenated Bitmaps and/or page-aligned blocks
- Cumulative Sum of precomputed rank values

What is a wavelet tree?


Wavelet Tree: Definitions

- In its basic form, the wavelet tree is a balanced binary tree.
- It stores a *sequence* $S[1, n] = c_1 c_2 c_3 \dots c_n$ of *symbols* $c_i \in \Sigma$, where $\Sigma = [1 \dots \sigma]$ is the *alphabet* of S .
- The tree has height $h = \lceil \log \sigma \rceil$, and $2\sigma - 1$ nodes, with σ of those as leaf nodes and $\sigma - 1$ as internal nodes.

Constructing the Wavelet Tree

- The wavelet tree is constructed recursively, starting at the root node and moving down the tree, with each node in the tree receiving a string constructed by its parent, except the root node that receives the full input string.
- Each node calculates the middle character of Σ and uses it to set the bits in the bitmap and split S in two substrings S_{left} and S_{right} .

Wavelet Tree Example



$S = \text{adsfadaadsfaads}$, $\Sigma = \text{adfs}$

Construction time and memory usage

- Construction time: $O(n \cdot h) = O(n \log \sigma)$
 - The Wavelet Tree can theoretically be constructed in $O(n \cdot h) = O(n \log \sigma)$ time as the sum of the lengths of the strings being processed at any single layer of the tree is the length of the input string to the tree.
- Memory usage: $O(n \log \sigma + \sigma \cdot ws)$ bits
 - At each level in the tree at most n bits are stored in the bitmaps in total, making $n \cdot h = n \cdot \log \sigma$ an upper bound to the total number of bits that a wavelet tree stores in its bitmaps.
 - In addition to this, each node takes some constant amount of machine words of space, and there are $2\sigma - 1$ nodes in the tree. ws is the size of our machine words. This makes the total memory consumption $O(n \log \sigma + \sigma \cdot ws)$ bits.

Queries


- The wavelet tree supports three queries:
 - **Access(p)**: Return the character c at position p in sequence S .
 - Running time: $O(n \log \sigma)$.
 - We have not implemented Access because it resembles Rank.
 - **Rank(c, p)**: Return the number of occurrences of character c in S up to position p .
 - Running time: $O(n \log \sigma)$.
 - **Select(c, o)**: Return the position of the o th occurrence of character c in S .
 - Running time: $O(n \log \sigma)$

Rank on a Wavelet Tree

$$\square = p$$

Query: Rank(c='a', p=10), Result: o = 3

0 = bit representing 'a'




Select on a Wavelet Tree

 = p

Query: $\text{Select}(c='a', o=3)$, Result: $p = 9$

0 = bit representing 'a'



Applications

- Information Retrieval


- Positional inverted index
- Document retrieval
- Range Quantile Query: Return the k th smallest number within a subsequence of a given sequence of elements.
- FM-count: Return number of occurrences of a pattern p in S .

- Compression
 - Zero-order entropy compression (H_0) using a RLE Wavelet Tree or a Huffman Shaped Wavelet Tree.
 - Higher-order entropy compression (H_k) using Burrows-Wheeler transformation and a RLE wavelet tree.
 - $H_k \leq H_0 \leq \log \sigma$.


Compression: Run-length encoding

- Run-length encoding counts the number of consecutive occurrences of a symbol and substitutes the consecutive occurrences with the symbol followed by its number of occurrences.
- Example: $RLE(\text{aaaaabbbaacccccaaaaa}) = \text{a5,b3,a2,c5,a5}$.
- Binary example: $RLE(00000000001111100000) = 10, 5, 5$
 - We can avoid specifying the symbol by assuming that 0 is always the first symbol.
 - If the binary number begins with a 1 we just add a 0 to the beginning of the result.
- Query by reversing RLE. It takes linear time $O(n)$ to reverse. Rank and select query time becomes $O(2n \log \sigma) = O(n \log \sigma)$
- Achieves space complexity within H_0

RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



(a) Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



(b) RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$

Compression: Burrows-Wheeler transform

- BWT permutes the order of the characters. If the original string had several substrings that occurred often, then the transformed string will have several places where a single character is repeated multiple times in a row.
- As a result it groups symbols more which improves the effect of Run-length encoding
- BWT is reversible
- Combined with RLE Wavelet Tree it achieves H_k compression.

BWT example

$S = \text{bananahat.}$

<i>bananahat#</i>	\Rightarrow	<i>#bananahat</i>
<i>ananahat#b</i>		<i>ahat#banan</i>
<i>nanahat#ba</i>		<i>anahat#ban</i>
<i>anahat#ban</i>		<i>ananahat#b</i>
<i>nahat#bana</i>		<i>at#bananah</i>
<i>ahat#banan</i>		<i>bananahat#</i>
<i>hat#banana</i>		<i>hat#banana</i>
<i>at#bananah</i>		<i>nahat#bana</i>
<i>t#bananaha</i>		<i>nanahat#ba</i>
<i>#bananahat</i>		<i>t#bananaha</i>

$BWT(S) = \text{tnnbhaaaa.}$

Burrows-Wheeler reverse transform example

$S = dca$

$$M = \begin{bmatrix} dca\# \\ ca\#d \\ a\#dc \\ \#dca \end{bmatrix} \Rightarrow M' = \begin{bmatrix} \#dca \\ a\#dc \\ ca\#d \\ dca\# \end{bmatrix}$$


$BWT(S) = acd$

Reverse BWT:


Add 1	Sort 1	Add 2	Sort 2	Add 3	Sort 3	Add 4	Sort 4
<i>a</i>	<i>#</i>	<i>a#</i>	<i>#d</i>	<i>a#d</i>	<i>#dc</i>	<i>a#dc</i>	<i>#dca</i>
<i>c</i>	<i>a</i>	<i>ca</i>	<i>a#</i>	<i>ca#</i>	<i>a#d</i>	<i>ca#d</i>	<i>a#dc</i>
<i>d</i>	<i>c</i>	<i>dc</i>	<i>ca</i>	<i>dca</i>	<i>ca#</i>	<i>dca#</i>	<i>ca#d</i>
<i>#</i>	<i>d</i>	<i>#d</i>	<i>dc</i>	<i>#dc</i>	<i>dca</i>	<i>#dca</i>	<i>dca#</i>

**#* = end of line character

RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$



(a) RLE Wavelet Tree on string *bananahat* with alphabet $\Sigma = abhnt$

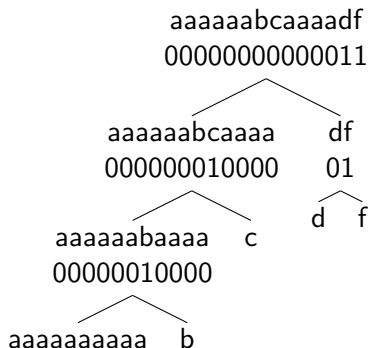


(b) BWT RLE Wavelet Tree on string *tnnbhaaaa* with alphabet $\Sigma = abhnt$


Huffman shaped wavelet tree

- Use Huffman codes of symbols to shape the tree
- A Huffman code is a binary value assigned to each symbol. The symbol with the highest frequency gets the lowest value.
- Shaping the tree based on Huffman codes places the most frequent symbols at the top of the tree and least frequent symbols at the bottom of the tree.
- Huffman shaping only makes sense on non-uniformly distributed data like a natural language text.

Huffman Shaped Wavelet Tree: Example



(a) Balanced Wavelet tree: 39 bits



(b) Huffman-shaped wavelet tree: 22 bits

Huffman Shaped WT: Space complexity

- Balanced version: $n \log \sigma + o(n \log \sigma) + O(\sigma \log n)$ bits
- Huffman-shaped: $n(H_0(S) + 1) + o(n(H_0(S) + 1)) + O(\sigma \log n)$ bits.
[Efficient Compressed Wavelet Trees over Large Alphabets by Navarro et al.]
- Huffman-shaped + Compressed Bitmap (RLE):
 $nH_0(S) + o(n(H_0(S) + 1)) + O(\sigma \log n)$ bits.


Experiments and Results

Focus of experiments


- Focus on optimizing and observing the effect of hardware penalties.
 - Cache Misses.
 - Branch Mispredictions.
 - Translation Lookaside Buffer (TLB) Misses.

- 1. Calculate binary rank and select using popcount
- 2. Pre-compute binary rank values in blocks
- 3. Block size dependence on input n
- 4. Pre-compute cumulative sums of rank values
- 5. Branchless select query
- 6. Queries on skewed cumulative sum wavelet tree

Calculate binary rank and select using popcount



(a) Rank




(b) Select


Figure : Rank and select queries using simple binary rank and select vs. rank and select queries using binary rank and select using the popcount instruction. Y-Axis is index 100 of the simple queries, that is, every value is percent of the value for the simple query.

Pre-compute binary rank values in blocks

- Pre-compute binary rank values in blocks to reduce amount of popcount calls.



(a) Rank



(b) Select

Figure : Comparison of wall time of rank and select queries between SimpleNaive not using precomputed values and UnalignedNaive using precomputed values.


Concatenated Bitmaps and/or page-aligned blocks

- Concatenated bitmaps
 - Save all bitmaps as one large bitmap to reduce memory usage by removing pointers to individual bitmaps.
 - Access using an offset (ulong) and a size for the bitmap (uint).
 - This uses $64 + 32 = 96$ bits per node vs. individual bitmaps with pointers taking up $3 \times 64 = 192$ bits per node.
 - Individual bitmaps for each node are word-aligned, and the bits between the last used bit and the end of the last used word are wasted.
- Page-aligned blocks
 - To reduce TLB misses
 - If the blocks are not aligned with the memory pages a block can be partly in one page and partly in another.
 - Page-aligning makes sure that blocks does not span page borders.

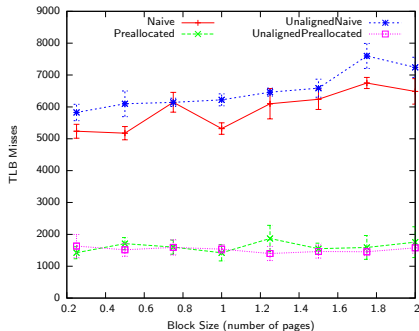
The various precomputed versions

Name	Concatenated Bitmaps	Page-aligned Blocks
Preallocated	yes	yes
UnalignedPreallocated	yes	no
Naive	no	yes
UnalignedNaive	no	no

Rank and select TLB misses




(a) Rank: TLB Misses




(b) Select: TLB Misses

- *Naive* does reduce TLB misses because of page alignment.
- Concatenated bitmaps reduces TLB misses, but page-aligning does not have much effect.

Running time: Pre-compute binary rank values in blocks



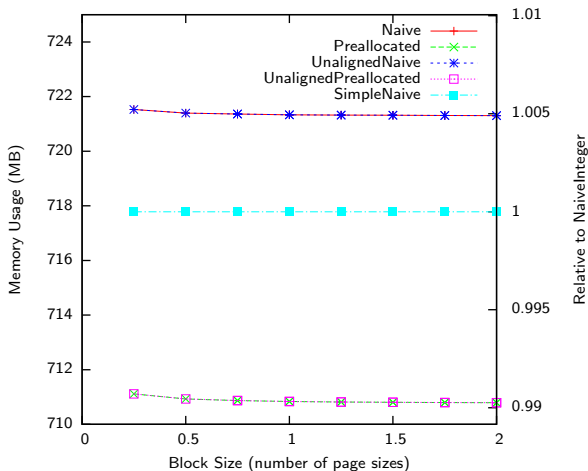
(a) Rank: Running Time



(b) Select: Running Time

Best Block size: $\frac{1}{2}$ page size = $\frac{1}{2} * 4096$ bytes = 2048 bytes.


Memory usage: Pre-compute binary rank values in blocks




Block size dependence on input size n

- When using lookups of precomputed values It costs $O(\frac{n}{b} + b)$ to calculate the binary rank.
- It costs $O(\frac{n}{b})$ to scan the blocks, and $O(b)$ to calculate the rank within a single block using popcount. The optimal block size should be one that minimizes this.
- The derivative of $\frac{n}{b} + b$ is $1 - \frac{n}{b^2}$ and its root is $n = b^2$ making the optimal block size $b = \sqrt{n}$.
- This is only the optimal block size for a single bitmap, and a wavelet tree has many bitmaps of varying sizes n that are lower near the leaves.



Experiment: Block size dependence on input size n for Rank



(a) $n = 10^5$



(b) $n = 10^6$



Cumulative sum



The End